

Problem → Build System for phony online & offline users

User 1 •
User 5 • (Not thinking about scaling)
User 4 •
User 3 •

→ Ways to Approach a System Design Problem

→ Spiral

decide the core &
Build around that
ex: Start Small Build
around it, get bigger

✓ we are using spiral
for the worst problem

Incremental

when we are unsure of how
to implement

- ① Start w/ dog zero archi
- ② See how each component
would behave
→ under load
→ at scale
- ③ Identify the bottleneck
- ④ Re-architect
- ⑤ Repeat

→ When the system is simple enough
→ go w/ Spiral

→ complex systems require incremental
→ startups work like this

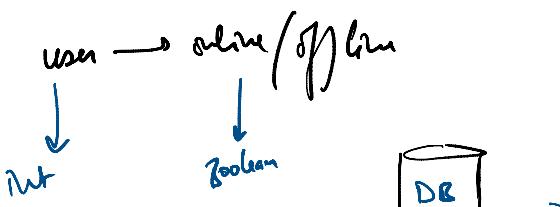


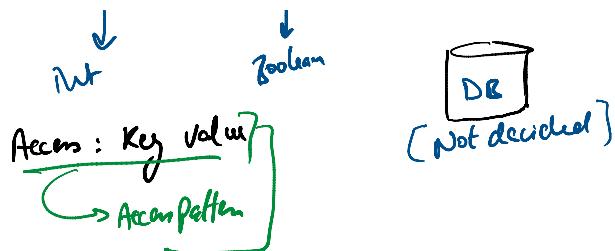
Remember

- understand the core property & access pattern
- don't stick to a particular tech
- aka don't be rigid
- ex: I'll always use PostgreSQL
I'll always use graph DB

→ Build an intuition towards
Building Systems

Storage





Why?
 we can fetch details for particular user
 so each id can be a key
 hence Key Value

Interfacing API



GET /players/Users?ids = u₁, u₂, u₃...

- design decision
 - it is better to have a endpoint that can handle Batch requests
 - individual request per user is a poor experience for the frontend

→ how to know if user is online



POST /heartbeat

↓ when we

we can decide per user case when we want to declare the user dead or the same.

like ⚡ No heartbeat in 30 sec, then Kill

↳ this means we need to store when we received the last heartbeat of every user

pulse		when you receive last heartbeat
user-id	last-hb	
u1	1000	UPDATE pulse
u2	1010	SET last-hb = now()
u3	1500	DELETE user-id = u1

get status API
 go on GET status becomes more complex
 GET /status / <user-id>
 → ⚡ No entry in DB → kill
 → ⚡ entry, but last-hb is > 30 sec
 ↓
 ⚡ kill

100 users	→ 100 entries	Each entry has 2 columns
1000 users	→ 1000 entries	
1m users	→ 1m entries	user-id , last-hb
1B users	→ 1B entries	↓ ↓
		int(48) int(48)
Total storage required for		Size of each entry = 8B
1B entries	= 8GB	

Can we do better on storage?
Requirement : User 1 • User 2 • User 3 • } all we care is online/ offline

what if absence == offline?

Idea: if user not present in the DB, we return offline.

So, let's expire the entries after 30 seconds
↓
delete

if we delete entries → we save a bunch of space
by not storing data of inactive users

Total entries = Active users

if 1B total users and 100k active then total entries = 100k
total size = 800 KB

How to auto delete?

Approach 1: Write a CRON job that deletes expired entries

- Not a robust solution
- We need to handle edge case in the business logic

Approach 2: Can we not offload this to our datastore?

Never re-invent the wheel!

DB with KV + expiration → Redis
Key value → DynamoDB }

Upon receiving an heartbeat

- update entry in Redis/ DynamoDB with ttl = 30 sec

Every heartbeat move the expiration time forward!

Which one would you pick & why? * cost estimation

arg

Which one would you pick & why? * cost estimation

Redis	DynamoDB
✓	✓ persistence
✓	managed (* item size) ? expensive
✓	version locking (cost/issue/ future consistency)
✓	time sensitivity

Always know your DB well
→ features
→ drawbacks etc

How is our DB doing?

Heartbeat is sent every 10 seconds

So, one user in 1 minute sends 6 heartbeats

If there are 1m active users, our system will get 6m req/min

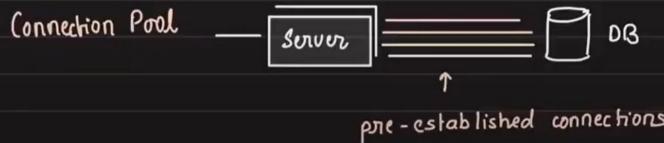
Each heartbeat request results in 1 DB call



How to make it better?

What's hectic? creating a new connection everytime

~~Very expensive~~



The microupdates that we were about to do for each online/offline event would require a new tcp connection to be established between server and DB, and then a teardown would be required (3 way handshake + 2 step teardown = 5 network calls for each update)

Setting up this connection is expensive, tearing it is also expensive.

Instead, we will use connection pool

Ex: we will setup some tcp connections with DB when server starts up, and maintain these connections. We will use these connections for the updates, and wont care about tear downs. This will help reduce setup+teardown, we will just use existing connections if available. This will reduce computation resources on both server and DB. Also, latency will be low as connection is pre-established

ISSUES

We need to set a minimum and maximum connection pool size.
Lets say minimum is 3. if 3 are occupied and one more comes
then two things can happen

- 1 - the new update waits till a connection is available
- 2 - a new connection is established

In case 2, the new connection established will not be terminated, but rather added to the connection pool, which basically means we are stretching the pool on demand.

We should limit this stretching, as the DB can only support so many tcp connections.

What if there are multiple servers, all with pooling? Will the DB be able to handle it?
For this, we should ensure unused connections are terminated after some time.

Remember, connection pooling is not a service, but a library/way of doing things.

We have a blocking queue + multiple connections.

We can imagine this as

```
Const connection1 = DB.connect()  
Const connection2 = DB.connect()  
...  
Queue.enqueue(connection1, connection2,...)
```

```
While(overCapacity) { wait() }  
Process(queue.top()); queue.dequeue();
```

^ tukka maara hai but probably this is what it looks like

Insight → we can use the same idea to check the health of systems in a distributed system
Similar pattern → observe the pattern

Foundation Topics in System Design

Database
Caching

Scaling
Delegation

Concurrency
Communication

Any and every decision would affect one of these 6 factors

We design a Multi-user blogging platform (Medium.com)

- one user multiple blogs
- multiple users

We take a look at key design decisions

Database

users	blogs
- id	- id
- name	- author_id
- bio	- title
	- body
	- is_deleted
	- published_at

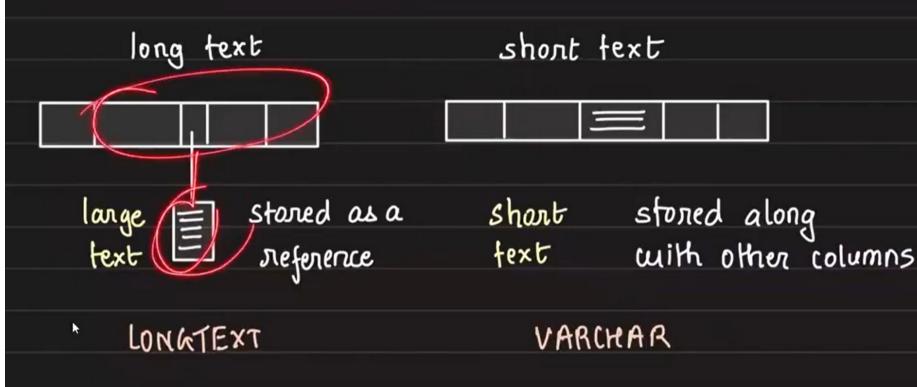
Importance of is_deleted [soft delete]

when user invokes delete blog, instead of DELETE we UPDATE

* Key reasons: Recoverability, Archival, Audit

+ Easy on the database engine [No tree re-balancing]

Column Type: body v/s bio



Storing datetime in DB	
datetime as datetime	02-04-2022T09:01:36Z
convenient sub-optimal heavy on size and index	serialized in some format
datetime as epoch integer	seconds since 1 st Jan 1970
efficient optimal, lightweight	172562349162
datetime as custom format (int)	YYYYMMDD - 20220402

takes 20B to store

| difference

takes 4B to store

Red Bus initially had this,
But ran into performance issues

didn't switch to epoch because they
wanted readability.

So they implemented their own thing
They only needed date, so they ignored
time & stored date as int

fast & smell

Caching

any

- Reduces response times by saving heavy computation

* Cache are not only RAM based

Typical use: reduce disk I/O or network I/O or compute

Caches are just glorified Hash Tables

+
with some advanced data structures

